



To make Medium work, we log user data. By using Medium, you agree to our [Privacy Policy](#), including cookie policy.

[Sign up](#)[Sign In](#)

# A Step-by-Step Guide to Creating Load Balancer and EC2 with Auto Scaling Group using Terraform

Achieving Scalability and High Availability using Load Balancer and Auto Scaling groups for your server.



Sharmila S · [Follow](#)

7 min read · Jun 19



2

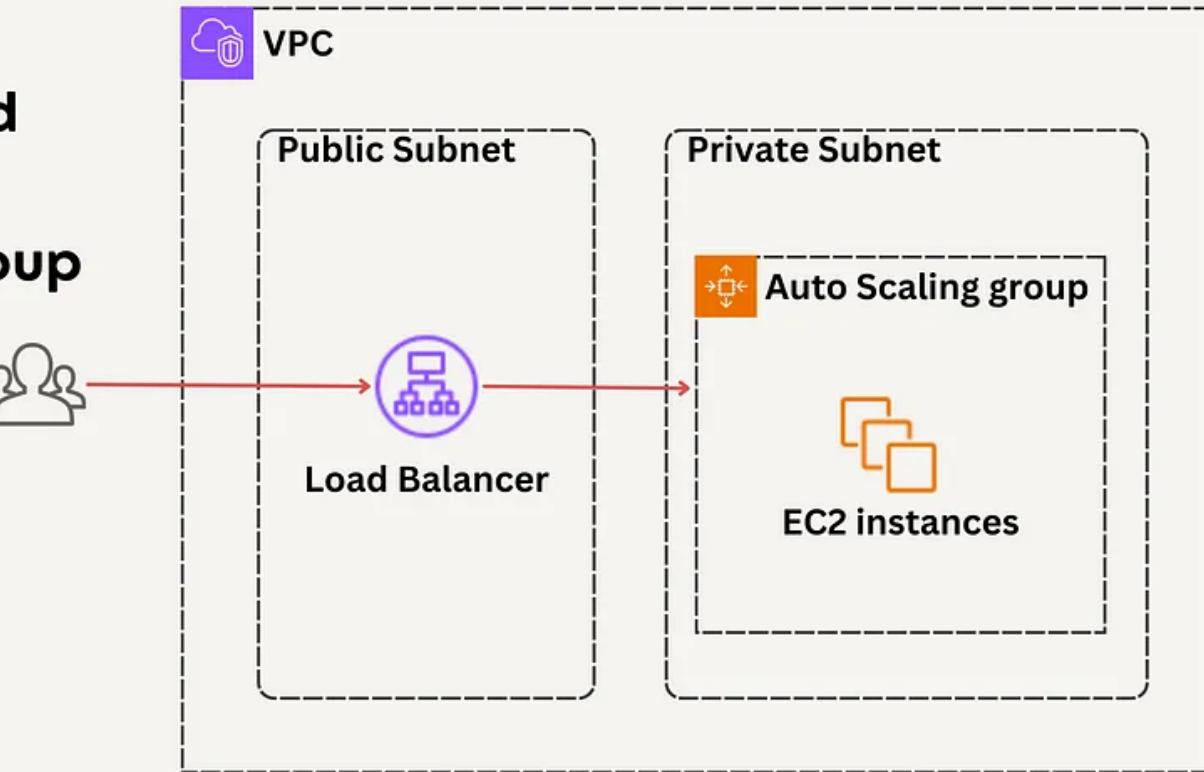


To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

# Deploy Application load balancer with Auto scaling group in AWS Cloud

with Terraform

@SharmilaS



**A**mazon Elastic Load Balancer (ELB) distributes incoming network traffic across multiple servers, preventing any single resource from

being overwhelmed. By evenly distributing the workload, load balancing

im To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

**A**mazon Auto Scaling Group (ASG) adjusts the number of EC2 instances based on the current demand. It provides elasticity, allowing the application to adapt to changes in traffic patterns. During periods of high demand, ASG provisions additional resources to handle the workload. During low demand, it removes excess resources and thereby reduces costs. It maintains high availability by automatically replacing failed servers with new instances.

**T**erraform is an infrastructure as code (Iac) tool that automates infrastructure provisioning. It is the most popular and offers support for multiple clouds.

## In this article, we will

1. Create a VPC, Subnets, and Configure network requirements.
2. Create EC2 instances with Auto scaling group and Launch template.
3. Create a Load balancer with a target group to the Auto Scaling group.

## Setting up the Environment

## 1. Install AWS CLI and configure

If you want to use this service, you must accept our Privacy Policy, including cookie policy.

### Installing or updating the latest version of the AWS CLI

Install the AWS CLI on your system.

[docs.aws.amazon.com](https://docs.aws.amazon.com/cli/latest/userguide/installing.html)

The next step is to configure your AWS profile to use in the CLI. Follow the instructions from the [official documentation](#) based on your machine type.

### Configuring the AWS CLI

Configure the AWS Command Line Interface (AWS CLI) and specify the settings for interacting with AWS.

[docs.aws.amazon.com](https://docs.aws.amazon.com/cli/latest/userguide/configure.html)

## 2. Install Terraform

Install Terraform in your system, if you don't have it already. Follow the instructions based on your platform on the official site given below.

[Install Terraform | Terraform | HashiCorp Developer](https://www.terraform.io/)

## Install Terraform on Mac, Linux, or Windows by downloading the

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

### 3. Initialise Project

Create an empty project directory and create a file called `main.tf` within it.

Firstly we need to add information about the provider (Here, we will be using AWS Provider). We will also configure AWS credentials, so we could provision resources in the AWS cloud.

```
1  terraform {  
2      required_providers {  
3          aws = {  
4              source  = "hashicorp/aws"  
5              version = "~> 4.0"  
6          }  
7      }  
8  }  
9  
10 # Configure the AWS Provider  
11 provider "aws" {  
12     region           = "us-east-1"  
13     shared_config_files = ["/Path/to/.aws/config"]  
14     shared_credentials_files = ["/Path/to/.aws/credentials"]  
15     profile           = "PROFILE"  
16 }
```

main.tf hosted with ❤️ by GitHub

[view raw](#)

You can configure the AWS credentials in the *provider* section. Here, I am using the path to AWS config and credential files in my system to authenticate. Also, I am mentioning the region and profile that I will be using by default. (You can also use Access and Secret keys instead of config files)

Now that we have set the provider set, we have to initialise Terraform to use it. (This command will download/update the providers).

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

**Initializing the backend...**

**Initializing provider plugins...**

- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v4.62.0

**Terraform has been successfully initialized!**

terraform init

## Networking requirements

We are going to create a VPC with 3 subnets (2 – public, 1 – private) across the region as,

- a public subnet in us-east-1a
- a public subnet in us-east-1b
- a private subnet in us-east-1b

*Note: You can create a new file for VPC configuration or continue in the same file.*

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

*id*

```
1 # VPC
2 resource "aws_vpc" "sh_main" {
3   cidr_block = "10.0.0.0/23" # 512 IPs
4   tags = {
5     Name = "sharmi-vpc"
6   }
7 }
8
9 # Creating 1st public subnet
10 resource "aws_subnet" "sh_subnet_1" {
11   vpc_id          = aws_vpc.sh_main.id
12   cidr_block      = "10.0.0.0/27" #32 IPs
13   map_public_ip_on_launch = true      # public subnet
14   availability_zone    = "us-east-1a"
15 }
16 # Creating 2nd public subnet
17 resource "aws_subnet" "sh_subnet_1a" {
18   vpc_id          = aws_vpc.sh_main.id
19   cidr_block      = "10.0.0.32/27" #32 IPs
20   map_public_ip_on_launch = true      # public subnet
21   availability_zone    = "us-east-1b"
22 }
23 # Creating 1st private subnet
24 resource "aws_subnet" "sh_subnet_2" {
25   vpc_id          = aws_vpc.sh_main.id
26   cidr_block      = "10.0.1.0/27" #32 IPs
```

```
27 map_public_ip_on_launch = false      # private subnet  
28  
29 To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy,  
including cookie policy.
```

v|

w|

## VPC and subnets

We can create a VPC using the resource `aws_vpc`. While creating VPC, it's important to choose a CIDR block that is large enough to accommodate the desired number of resources, but not too large to waste IP addresses. (Here we have assigned a pool of 512 IP addresses for this tutorial)

We create subnets using the `aws_subnet` resource by providing,

- the VPC id
- CIDR block of IP addresses (Make sure the IP address block is within the IP address range of the VPC)
- and the availability zone.

For the public subnet, we have to provide `map_public_ip_on_launch` as *true* (Default — *False*).

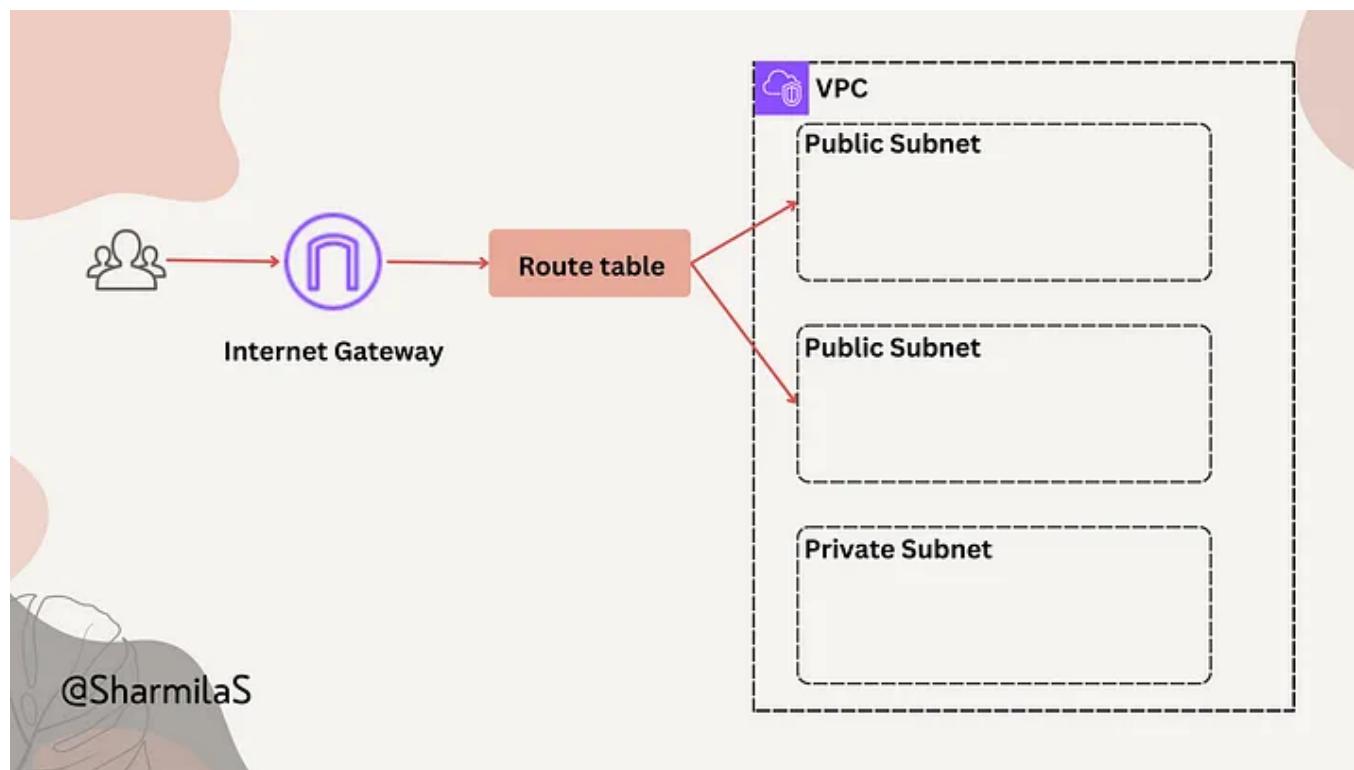
## Route tables and Gateways

## 1. Gateway for public subnets

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

Subnets are the smallest unit of a VPC. They define the IP range and routing rules for traffic within a specific part of the VPC.

We can create a route table for our public subnet to connect it to Internet Gateway.



Let's create an Internet Gateway and a route table. Then, associate this route table] To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

```
1 # Internet Gateway
2 resource "aws_internet_gateway" "sh_gw" {
3   vpc_id = aws_vpc.sh_main.id
4 }
5
6 # route table for public subnet - connecting to Internet gateway
7 resource "aws_route_table" "sh_rt_public" {
8   vpc_id = aws_vpc.sh_main.id
9
10  route {
11    cidr_block = "0.0.0.0/0"
12    gateway_id = aws_internet_gateway.sh_gw.id
13  }
14 }
15
16 # associate the route table with public subnet 1
17 resource "aws_route_table_association" "sh_rta1" {
18   subnet_id      = aws_subnet.sh_subnet_1.id
19   route_table_id = aws_route_table.sh_rt_public.id
20 }
21 # associate the route table with public subnet 2
22 resource "aws_route_table_association" "sh_rta2" {
23   subnet_id      = aws_subnet.sh_subnet_1a.id
24   route_table_id = aws_route_table.sh_rt_public.id
25 }
```

gateways-public.tf hosted with ❤️ by GitHub

[view raw](#)

## Internet gateway for the public subnets

To make Medium work, we log user data. By using Medium, you agree to our [Privacy Policy](#), including cookie policy.

No, we can't do that. We can't associate a route table to a subnet.

## 2. Gateway for Private Subnet

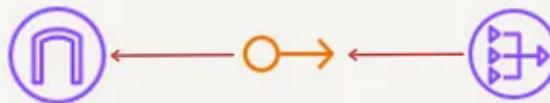
We will be creating our EC2 instances in the private subnet, allowing only the requests from the load balancer to reach the instances.

But, the instances might need to connect to the internet to download software/tools. To provide access to the internet, we need a NAT gateway for the private subnet.

*The NAT instance must have internet access, So, it must be in a public subnet (a subnet that has a route table with a route to the internet gateway), and it must have a public IP address or an Elastic IP address.*

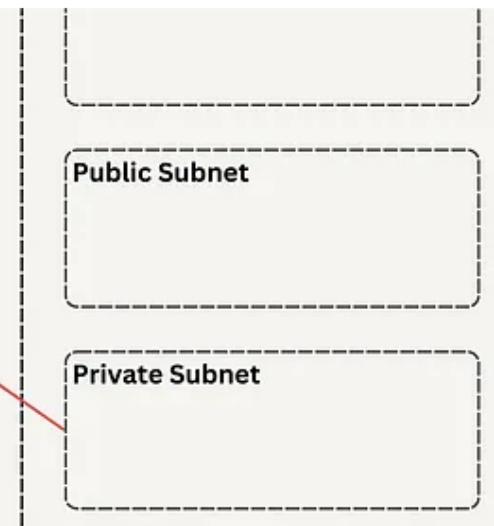
To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

Internet Gateway      NAT Gateway



Elastic IP

Route table



Public Subnet

Private Subnet

@SharmilaS

Creating a route table and associating it with the subnet is the same process as the previous one.

```
1 # Elastic IP for NAT gateway
2 resource "aws_eip" "sh_eip" {
3   depends_on = [aws_internet_gateway.sh_gw]
4   vpc        = true
5   tags       = {
6     Name = "sh_EIP_for_NAT"
7   }
8 }
```

```
9  
10  
11 To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy,  
12 including cookie policy.  
13  
14   allocation_id = aws_subnet.sh_subnet_1.id # nat should be in public subnet  
15  
16   tags = {  
17     Name = "Sh NAT for private subnet"  
18   }  
19  
20   depends_on = [aws_internet_gateway.sh_gw]  
21 }  
22  
23 # route table - connecting to NAT  
24 resource "aws_route_table" "sh_rt_private" {  
25   vpc_id = aws_vpc.sh_main.id  
26  
27   route {  
28     cidr_block      = "0.0.0.0/0"  
29     nat_gateway_id = aws_nat_gateway.sh_nat_for_private_subnet.id  
30   }  
31 }  
32  
33 # associate the route table with private subnet  
34 resource "aws_route_table_association" "sh_rta3" {  
35   subnet_id      = aws_subnet.sh_subnet_2.id  
36   route_table_id = aws_route_table.sh_rt_private.id  
37 }
```

gateways-private.tf hosted with ❤ by GitHub

[view raw](#)

NAT gateway for the private subnet

## Configuring the Load Balancer

We To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

handling traffic and hitting requests.

Remember to provide a public subnet(s) for our load balancer, since we need it to be internet-facing.

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

*Let's proceed to Auto Scaling group configuration first and cover the security,*

*gr* To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

## Creating the Auto Scaling Group

We will be creating and using an EC2 Launch Template to create EC2 instances.

*You can also provide the EC2 configurations directly in the aws\_autoscaling\_group resource instead of creating a launch template to use*

In the launch template, we are providing,

- the AMI we are going to use (make sure to get the AMI id from the same region as the EC2 instance)
- the instance type
- user data script (base64 encoded file which has commands to start *Apache web server*)
- security group, subnet id for network configurations.

*We will create the security group mentioned here in the next section*

For the Auto scaling group,

## 1 Configuring the capacity

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

- `-` `capacity` specifies the maximum number of instances that the scaling group (ASG) can have, setting a limit to prevent excessive resource usage or costs.
  - `min_size` sets the minimum number of instances that the group should always have, ensuring a baseline capacity during low demand.
  - `desired_capacity` is the desired number of instances that the ASG should aim to have at any given time. It represents the ideal number of instances to handle the current workload.
2. We can connect this Auto Scaling Group to the Load Balancer target group using `target_group_arns`.
3. We add the launch template we created and also the private subnet.

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

Auto Scaling group with Launch template

## **My user data file: (to run Web server)**

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

```
sudo systemctl start httpd
sudo systemctl enable httpd
echo "<h1>Hello World from $(hostname -f)</h1>" > /var/www/html/index.html
```

## Security access rules — for EC2 and Load Balancer

Let's create a security group each for load balancer and the EC2 instances.

For the Load Balancer, we are allowing HTTP and HTTPS requests from the internet.

For the EC2 instances, we allow only HTTP requests from the load balancer.

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

Security groups for ELB and EC2

## Terraform apply

We have all the resources ready!

Let's look at the terraform commands that we would be needing every time we deploy any changes.

To make Medium work, we log user data. By using Medium, you agree to our [Privacy Policy](#), including cookie policy.

## Terraform Important Commands

<b>terraform fmt</b>	Formats the Tf files
<b>terraform validate</b>	Validates the resources and configurations offline
<b>terraform plan</b>	Gives a view of changes to be deployed
<b>terraform apply</b>	Deploy the resources to the Cloud

@SharmilaS

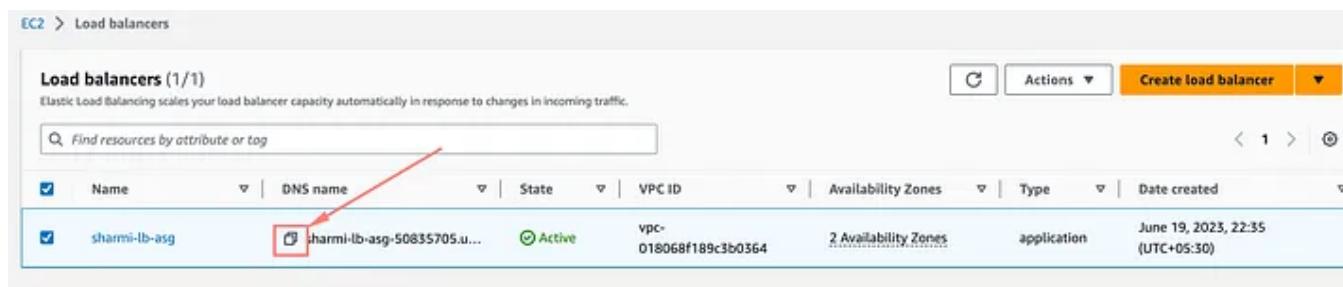
Try the above commands in that order.

~~terraform plan gives you details about the resources that will be deployed~~

an To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.  
de

After the changes are deployed,

- Go to [AWS Console](#) and sign in.
- Go to Load Balancer under EC2 service.
- Select the load balancer that we created and copy the DNS name.



The screenshot shows the AWS EC2 Load Balancers page. The title bar says "EC2 > Load balancers". Below it, a header row includes "Load balancers (1/1)", a search bar, and buttons for "Actions" and "Create load balancer". A red arrow points from the text "Select the load balancer that we created and copy the DNS name." to the "DNS name" column in the table below. The table has columns: Name, DNS name, State, VPC ID, Availability Zones, Type, and Date created. One row is visible: "sharmi-lb-asg" (selected), "sharmi-lb-asg-50835705.u...", "Active", "vpc-018068f189c3b0364", "2 Availability Zones", "application", "June 19, 2023, 22:35 (UTC+05:30)".

Name	DNS name	State	VPC ID	Availability Zones	Type	Date created
sharmi-lb-asg	sharmi-lb-asg-50835705.u...	Active	vpc-018068f189c3b0364	2 Availability Zones	application	June 19, 2023, 22:35 (UTC+05:30)

- Hit the URL in the browser

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

# Hello World from ip-10-0-1-9.ec2.internal

You can make changes in the infrastructure and try the above commands in the same order, to deploy them. The terraform state files save the current state of the infrastructure.

## **Backup:**

The terraform state files (`terraform.tfstate` and `terraform.tfstate.backup`) can be backed up in storage — using which we can provision the resources again anytime if something goes wrong.

**Note:** It is recommended to make changes in the terraform code once

de To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy,  
including cookie policy.

tr~~e~~

manually, it can lead to inconsistencies between the desired state in your Terraform code and the actual state of the resources. This can cause conflicts and make it difficult to manage and update your infrastructure accurately.

## Terraform destroy

If you want to destroy all the resources created using Terraform in this project, you can use `terraform destroy` command.

This could be particularly helpful when you are learning Terraform, you can delete all the resources easily at the end, to avoid costs.

If you have any doubts, or any suggestions to improve this article, feel free to add them in the comments.

## GitHub Repo

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

### for creating a load balancer and auto...

Terraform code for creating a load balancer and auto scaling group.

- GitHub - SharmilaS22/medium-tf-ec2-alb-asg...

[github.com](https://github.com/SharmilaS22/medium-tf-ec2-alb-asg)

Hope this article helps! Follow for more!!

AWS

Aws Ec2

Cloud Computing

Technology

DevOps



## Written by **Sharmila S**



273 To make Medium work, we log user data. By using Medium, you agree to our [Privacy Policy](#), including cookie policy.

Sofware Engineer, [Writer about cloud & devops](#), [http://sharmilas.bio.link/](#)

### More from Sharmila S



 Sharmila S in featurepreneur

## Connect MongoDB database to Express server—step-by-step

4 min read · May 2, 2021

104

3



238

1



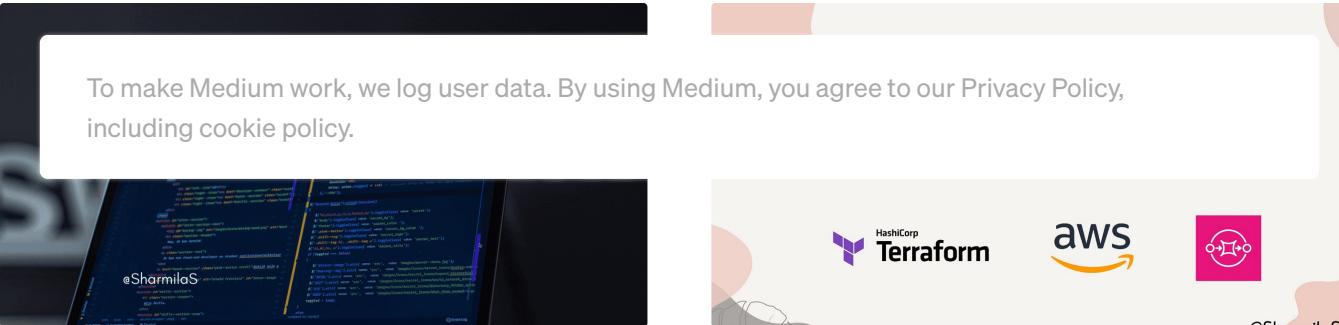
 Sharmila S

## Get Started with RabbitMQ in Node.js

In this article, we are going to connect two node.js applications with a queue using...

6 min read · Jan 28, 2022



 Sharmila S

## Download files for the client from the Node.js server using this quic...

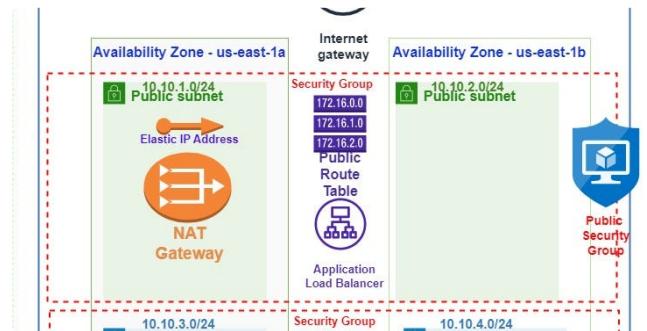
Add a download option for files on your Node.js web application for users to...

3 min read · Nov 30, 2021

 72 9[See all from Sharmila S](#)

To make Medium work, we log user data. By using Medium, you agree to our Privacy Policy, including cookie policy.

## Recommended from Medium



Jason Li ⚡ in AWS in Plain English

### How to create an AWS EC2 Auto-Scaling Group for High Availability

Welcome to week 21 with LUIT! As we continue to explore and learn the Efficiency,...

12 min read · Jun 3

👏 140



B8

### Terraform manage multiple environments

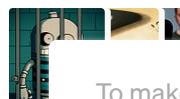
How to manage TF multiple environments in your projects

13 min read · Jul 24

👏 67



## Lists

**AI Regulation****ChatGPT prompts**

To make Medium work, we log user data. By using Medium, you agree to our [Privacy Policy](#), including cookie policy.

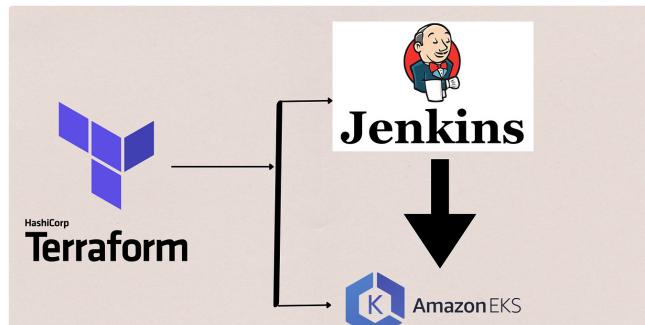


21 stories · 210 saves

ed

**Reading**

52 stories · 325 saves



Mubin Khalife

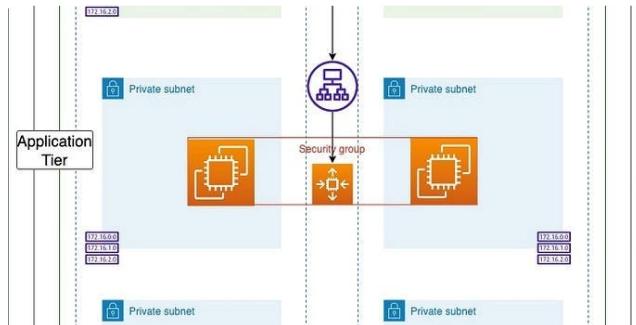
## DevOps project using Terraform, Jenkins, and EKS

In this article, we are going to create a DevOps project where we'll use Terraform,...

12 min read · Jun 20

22

1



priyanka kumari

## How to build 3 -tier architecture in AWS

What is a 3 Tier Architecture

12 min read · Aug 4

8





To make Medium work, we log user data. By using Medium, you agree to our [Privacy Policy](#), including cookie policy.



Dhruvin Soni

## How to import existing AWS resources using Terraform?

How to import existing AWS resources using Terraform?

6 min read · Apr 24

👏 13

Q 1

+W

Navidehbaghaifar

## Designing an AWS Three-Tier Architecture

In this week's project, I will guide you through the steps of constructing an AWS Three-Tier...

17 min read · Jun 23

👏 6

Q 1

+W

See more recommendations